

Hochschule Karlsruhe
Winter Semester 2025/2026
Course: Business Ethics
Supervising Professor: Prof. Dr. Bernd Banke

Term Paper

The Case Study: Deliveroo's Frank Algorithm

Ethical Risk of Strategic Reclassification of AI Models

Submitted by	: Nivedhidha Ilangovan
Matriculation No.	: 193522
Study Program	: Enterprise and IT Security

Table of Contents

Introduction.....	3
Overview of the EU Artificial Intelligence Act.....	3
The High-Risk Classification and the Article 6(3) Derogation.....	3
Problem Statement: The Ethical Risk of Strategic Reclassification.....	4
Awareness.....	4
Case Study Background: Deliveroo's Frank Algorithm.....	4
Theoretical Framework: Social Systems and Power Asymmetry.....	6
Structural Coupling: Economic System vs. Legal System.....	7
Analysis.....	8
Identification of Primary and Secondary Stakeholders.....	8
Primary Stakeholders.....	8
Secondary Stakeholders.....	8
The Mechanism of Information Asymmetry and the Black Box Effect.....	9
Strategic Reclassification as an Ethical Failure.....	9
The Motivation for Strategic Reclassification.....	10
The Nature of the Ethical Failure.....	10
The Violation of the Categorization Principle.....	10
Analysis.....	11
Utilitarianism.....	11
Deontological Ethics.....	11
Ethics of Responsibility.....	12
Ethics of Institutions.....	12
The Proposed Course of Action.....	12
Application.....	13
The Weighted Representation Model.....	13
The Ethics of Disclosure and Consensus.....	13
Continuous Monitoring and Feedback Loops.....	14
Conclusion.....	14
References.....	15

Introduction

Overview of the EU Artificial Intelligence Act

The European Union Artificial Intelligence Act (EU AI Act) represents the world's first comprehensive horizontal legal framework for the regulation of artificial intelligence. Established in 2024, the Act seeks to balance the promotion of technological innovation with the protection of fundamental human rights, safety, and democratic values. At its core, the legislation is built upon a risk-based approach, which assigns regulatory obligations to AI providers and deployers based on the level of potential harm a system poses to society.

The Act categorizes AI systems into four distinct levels of risk:

1. **Unacceptable Risk:** AI systems that pose a clear threat to safety or fundamental rights are strictly prohibited. (Example: social scoring by governments)
2. **High-Risk:** Systems that significantly impact an individual's life chances or physical safety (such as AI used in healthcare, critical infrastructure, or employment management) are permitted but are subject to rigorous compliance requirements, including data quality standards, human oversight, technical documentation, and quality management systems.
3. **Limited Risk:** Systems like chatbots or deepfakes are subject to lighter transparency obligations, ensuring users are aware they are interacting with AI.
4. **Minimal or No Risk:** The vast majority of AI applications, such as spam filters or AI-enabled video games, are largely unregulated.

The framework shifts the responsibility of safety from a reactive model to a proactive one. Companies must now conduct a **Conformity Assessment** before a product even enters the market. However, the flexibility within these risk classifications—particularly the ability for companies to self-classify certain systems as **non-high-risk**—creates a significant ethical dilemma regarding corporate transparency and accountability.

The High-Risk Classification and the Article 6(3) Derogation

The **High-Risk** category is the most operationally significant tier of the EU AI Act, as it encompasses systems with the potential to cause significant harm to health, safety, or fundamental rights. Under **Article 6(2)**, any AI system falling into the categories listed in **Annex III** is automatically considered high-risk. These categories include sensitive areas such as biometrics, critical infrastructure, and, most relevant to this study, employment and worker management. Systems in this tier must comply with strict mandates for data governance, human oversight, and cybersecurity before they can be legally deployed in the EU market.

But as stipulated in **Article 6(3) and 6(4) of the EU AI Act**, providers of systems listed in Annex III may derogate from high-risk obligations if they conclude the system does not pose a significant risk of harm or materially influence the outcome of decision-making. This applies if the AI performs purely preparatory tasks, improves human activity, or executes a narrow procedural task. However, the Act clarifies that if the system performs profiling of natural persons, it is always considered high-risk, regardless of these exceptions.

This self-assessment mechanism, while intended to prevent over-regulation of harmless tools, places the burden of proof on the provider. Under **Article 6(4)**, a provider opting for this derogation must document their assessment and register the system in an EU database, making it available to national authorities upon request. From a GRC perspective, this creates a documentation gap where the ethical integrity of the system depends entirely on the provider's honesty in defining what constitutes a narrow task versus a life-altering decision.

Problem Statement: The Ethical Risk of Strategic Reclassification

The ability for providers to self-classify **Annex III** systems as non-high-risk under **Article 6(3)** introduces a significant ethical challenge. Because high-risk compliance requires substantial financial investment and rigorous auditing, companies face a conflict of interest between short-term profitability and long-term safety. As noted in research on algorithmic management, firms often target growth by treating human workers as resources that can easily be substituted.

If a provider uses technical jargon to frame a high-impact system as a narrow task, they effectively bypass critical protections for employees and customers. This creates an information asymmetry where the business takes unfair advantage of the users, who remain unaware of how the black box affects their rights. This lack of responsibility undermines the primary goal of the EU AI Act: creating a trustworthy AI ecosystem.

Awareness

Case Study Background: Deliveroo's Frank Algorithm

Deliveroo is a prominent platform-based technology company operating in the gig economy, providing food delivery services by connecting customers, restaurants, and a fleet of self-employed riders. To manage this massive operation at scale, the company relies heavily on algorithmic management to coordinate its workforce. Central to this business model is an algorithm named Frank, which Deliveroo utilized to manage rider performance and determine access to work. In the gig economy context, such algorithms are essential for infinite scalability, as they minimize the need for interpersonal interaction between management and workers, essentially treating human labor as a calculable resource.

The Frank algorithm functioned primarily through a booking system where riders were required to reserve delivery windows each Monday for the upcoming week. Access to these windows was split into three prioritized time slots: 11:00 am, 3:00 pm, and 5:00 pm. Because delivery opportunities are finite, the 11:00 am slot was the most desirable, offering the highest chance of securing a rider's preferred shifts. However, access to these slots was not equal; it was determined by a reputation ranking calculated by the algorithm. This ranking was based on two specific variables: a peak participation index, which tracked the number of high-demand shifts a rider worked, and a reliability index, which was inversely proportional to the number of late cancellations. A cancellation was considered late if it occurred less than 24 hours before a booked shift or if a rider failed to log into the app at least 15 minutes before their window started.

The ethics of this system were challenged in the Court of Bologna by Italian trade unions, who alleged that the Frank algorithm was inherently discriminatory. The court's investigation centered on how the algorithm allocated work and whether its logic infringed upon the fundamental rights of the riders. The primary finding was that the algorithm was blind to the reasons behind a rider's behavior. Whether a rider cancelled a shift due to a strike, a serious illness, or a childcare emergency, the algorithm treated the event identically to a rider who simply chose not to work. By failing to differentiate between legitimate reasons for absence and mere unreliability, the algorithm effectively penalized workers for exercising their legal rights or attending to human necessities.

The legal and ethical discussion highlighted a significant information asymmetry, where the company held the black box logic of the algorithm while the riders suffered the consequences of its blind rankings. The court eventually ruled that the algorithm possessed an intrinsic discriminatory character. Significantly, the court noted that Deliveroo had the technical capability to prevent this discrimination—as the algorithm already accounted for exceptions like app crashes or consecutive shift injuries—but the company had deliberately chosen not to include broader legitimate reasons for absence.

As a result, the Court issued a landmark ruling, ordering Deliveroo to pay €50,000 in damages to the trade unions and to publish the court's decision on its website to ensure transparency. While Deliveroo argued that no objective and real discrimination had been proven, the company's operational response was significant. Rather than attempting to fix the existing code to include complex human exceptions, Deliveroo ultimately decommissioned the specific Self-Service Booking (SSB) system in the Italian market, transitioning toward a Free Login model to mitigate further legal liability. This shift was not just a return to business as usual; it was a transition into a more regulated era where black box logic in employment is no longer legally defensible. This remains highly relevant today as a textbook example of why employment-related AI is classified as High-Risk under **Annex III of the EU AI Act**, emphasizing that systems affecting human livelihoods must be subject to rigorous transparency and human oversight.

Theoretical Framework: Social Systems and Power Asymmetry

To understand why the Deliveroo case represents an ethical dilemma rather than just a technical error, it is useful to analyze it through the lens of Social Systems Theory (Luhmann's framework). Society is composed of different functional systems, each operating under a specific binary code. In this case, three primary systems—Economic, Legal, and Regulatory—interact, often with conflicting priorities.

- **The Economic System (Code: Profit / Loss):** The function of this system is to ensure the scarcity of resources is managed efficiently to ensure the survival of the firm. For Deliveroo, the Frank algorithm was a masterpiece of economic logic. By reducing human workers to data points, the company maximized Reliability and Peak Participation, directly translating to lower operational costs and higher scalability. In this system, if an action increases profit, it is coded as positive; if it incurs costs (like the high price of EU AI Act compliance) without a visible return, it is coded as a loss and avoided.
- **The Legal System (Code: Legal / Illegal):** The legal system's function is to provide normative certainty—ensuring that fundamental rights and laws are upheld consistently. While the Economic system saw an efficient tool, the Legal system (represented by the Bologna Court) saw a violation of labor protections. The binary here is simple: either the algorithm respects the rights of the workers (legal), or it discriminates against them (illegal). The court's ruling against Deliveroo was a corrective measure, asserting that economic efficiency cannot exist outside the boundaries of the law.
- **The Regulatory System (Code: Compliant / Non-Compliant):** Regulatory bodies act as the bridge between law and technology. Their function is to maintain order through standardized rules. As noted in the EU AI Act dilemma, if a company provides documentation that satisfies the technical requirements of **Article 6(3)**, the system is coded as Compliant. The ethical risk lies in the flexibility of the EU AI Act; while designed to encourage AI growth, this flexibility allows companies to hide behind technical jargon to achieve a Compliant status in paperwork, even when the reality of the system remains illegal in the eyes of the law.

The intersection of these systems reveals a profound Power Asymmetry. The business system holds the Information Power—it designs the algorithm and chooses how to present it to the regulatory system. The workers, meanwhile, are subjects of the Economic system's efficiency but are often invisible to the Regulatory system until a Legal intervention occurs. This gap between paper compliance and human reality is the primary ethical challenge for modern GRC.

Structural Coupling: Economic System vs. Legal System

The provided matrix illustrates the structural coupling between the economic and legal systems. Structural coupling occurs when a system's internal logic is irritated or influenced by its environment—in this case, the legal verdict forcing a change in economic operations.



Fig. 1: Structural Coupling between the Economic and Legal Systems

The trajectory of the Frank algorithm can be analyzed through four distinct phases of coupling:

- **Initial State (Profit / Illegal):** Before the court's intervention, Deliveroo operated based on a purely economic logic of infinite scalability and efficiency. The system was blind to external legal norms regarding worker rights, prioritizing the profit code while operating in a state of latent illegality due to its discriminatory nature.
- **The Verdict (Loss / Illegal):** The verdict from the Court of Bologna served as a massive disruption to the economic system. By imposing a €50,000 fine and requiring public disclosure of the ruling, the legal system forced the company to recognize that its profit-seeking logic had breached the boundaries of the law.
- **Operational Transition (Loss / Legal):** To avoid further legal liability, the company moved toward the Legal side of the binary by deactivating the specific ranking system. While this initially incurred a cost to operational efficiency—moving the company toward the Loss quadrant—it was a necessary step to align with regulatory expectations.

- **Adaptation (Profit / Legal):** The final adaptation into a Free Login model represents a successful structural coupling. The company found a new way to operate profitably while remaining within the boundaries of the Legal code. This ensured that the economic innovation respects the fundamental rights mandated by the legal system.

Analysis

Identification of Primary and Secondary Stakeholders

To analyze the ethical implications of the Frank algorithm, we must identify the actors within the system's reach. Stakeholders are categorized based on the directness of their relationship with the AI's automated decisions.

Primary Stakeholders

Primary stakeholders are those who are directly involved in the system's operation or whose fundamental rights are immediately impacted by its outputs.

- **The Riders (Data Subjects):** As the data points of the algorithm, their livelihood and ability to access work depended entirely on the Reliability and Peak Participation indices. They are the primary victims of the discrimination found by the court.
- **Deliveroo (The Provider/Deployer):** The company designed and deployed the algorithm to achieve infinite scalability and economic efficiency. They hold the power over the code and the data it processes.

Secondary Stakeholders

Secondary stakeholders are individuals or groups who influence the environment, set the rules, or intervene when a conflict arises, though they are not related directly.

- **Trade Unions:** In this case, the unions acted as essential intermediaries, representing the collective interests of the riders and challenging the black box logic that individual workers could not contest alone.
- **The Judiciary (The Courts):** The Court of Bologna acted as the moral and legal auditor, determining that the algorithm's logic was discriminatory and enforcing a penalty.
- **Regulatory Bodies (Future EU AI Enforcers):** While the EU AI Act was not in effect during the Deliveroo case, the case itself served as a pre-regulatory catalyst. These bodies now use such cases to define High-Risk categories in **Annex III** to ensure future companies cannot use the same blind logic.

The Mechanism of Information Asymmetry and the Black Box Effect

A critical issue in the Deliveroo case was the profound information asymmetry between the platform and its riders. While Deliveroo provided basic information about how the Reliability and Peak Participation indices affected shift access, the actual backend logic—how these variables were weighted and processed—remained a Black Box. Without access to the source code or a clear explanation of the algorithm's decision-making process, it was impossible for individual workers to verify if the system was operating fairly or if it was infringing upon their rights.

This lack of transparency led to several interconnected failures:

- Because the algorithm's internal logic was hidden, the fact that it failed to respect fundamental labor rights—such as the right to strike or sick leave—remained unknown until the court's intervention.
- For riders who survive solely on this salary, the blindness of the algorithm is a moral crisis; being denied a living because of a childcare emergency or illness is inherently unfair.
- The system did not just face an ethical dilemma; it suffered from a dual failure of both ethics and law. The Black Box allowed the company to bypass legal protections by framing a life-altering decision-making tool as a simple procedural ranking.
- An individual worker cannot argue against a mathematical output they do not understand. This Information Power held by the company effectively silenced workers until the trade unions provided the collective power to demand an audit.

This specific mechanism of hiding high-impact logic inside a Black Box is precisely why the EU AI Act now mandates strict transparency and human oversight for employment-related AI. Under **Annex III**, providers can no longer claim a system is low risk if its backend logic materially influences a person's access to work and livelihood.

Strategic Reclassification as an Ethical Failure

Strategic reclassification occurs when a provider deliberately categorizes an AI system into a lower risk tier to bypass the rigorous obligations of the EU AI Act. In the context of Deliveroo, the Frank algorithm was treated as a simple procedural tool for efficiency, despite the fact that it performed high-stakes management functions that dictated the livelihoods of thousands of riders.

The Motivation for Strategic Reclassification

The company's decision to rely on an automated, blind algorithm was driven by several strategic goals:

- By automating worker management, the company could minimize the need for HR personnel and interpersonal management, allowing for infinite scalability.
- Avoiding the High-Risk designation allows a firm to escape the financial and administrative costs of mandatory data auditing, human oversight, and cybersecurity protocols.
- If a system is classified as low-risk, the provider avoids the burden of proof regarding fairness and non-discrimination.

The Nature of the Ethical Failure

The ethical failure of the company leadership falls into one of two categories, both of which are indefensible:

- **Ethical Ignorance (Incompetence):** If the leadership truly did not realize the algorithm was discriminatory, they failed in their fundamental duty to understand the tools governing their workforce. A competent leader is responsible for the impact of their technology on all stakeholders.
- **Intentional Negligence (Malice):** If the company knew the Black Box was operating illegally but chose to maintain it to preserve profit, they actively participated in an ethical breach. The Clifford Chance report confirms that Deliveroo had the technical capability to build exceptions into the code but chose not to.

The Violation of the Categorization Principle

Classifications in the EU AI Act are not suggestions; they are safeguards designed to protect human rights. By attempting to use loopholes to change the category, the company prioritized short-term control over the safety of its employees.

True leadership requires balancing the Economic System (profit) with the Legal System (rights). An executive participating in decision-making must consider that if they would not want their own careers governed by a blind and unexplainable algorithm, they should not impose such a system on their workers. Using reclassification to escape accountability is a failure to protect the very stakeholders—the workers—who make the company's profit possible.

Analysis

Utilitarianism

Utilitarianism is a consequentialist framework where the morality of an action is determined by its ability to produce the greatest good for the greatest number or to minimize total suffering.

In the context of the Frank algorithm, a narrow utilitarian approach might initially favor the company's decision to perform strategic reclassification. By labeling the system as low-risk, the company reduces compliance costs and maintains infinite scalability, which theoretically benefits customers through lower prices and the company through higher profits.

However, a comprehensive utilitarian analysis reveals a net loss in societal utility. The black box logic caused widespread suffering for thousands of riders who faced financial instability and the erosion of their fundamental rights. Furthermore, the eventual legal consequences, including the €50,000 fine and the resulting reputational damage, represent a significant long-term loss that outweighs the short-term gains of avoiding regulation.

Therefore, a sustainable utilitarian approach would conclude that accurate classification and transparency are more beneficial, as they foster long-term trust and prevent the sufferings caused by litigation and labor unrest.

Deontological Ethics

Deontological ethics states that the morality of an action is based on whether it adheres to universal rules and duties, regardless of the consequences. A core principle here is that one should act only according to maxims that they would want to become universal laws.

When Deliveroo executives utilized black box rankings and strategic reclassification to manage workers, they violated the duty to treat individuals as ends in themselves rather than mere calculable resources. If every company were to universally adopt a policy of hiding life-altering logic to escape accountability, the entire labor market would collapse into a state of manipulation and mistrust. No executive would want their own career or livelihood governed by an unexplainable, blind algorithm that ignores human necessity.

From a deontological perspective, accurate classification is a moral duty because it ensures fairness and respects the autonomy of the worker, fulfilling the requirement that business practices must be capable of being practiced as a universal law of justice.

Ethics of Responsibility

The Ethics of Responsibility shifts the focus from the internal motives of the decision-maker to the actual consequences of their actions in the world. This model is particularly relevant for the unintentional executive who may not fully understand the technical complexities of AI.

Under this framework, a leader cannot claim ignorance as an excuse; they are responsible for the tools they deploy. Even if the Deliveroo executives intended to foster company growth or believed the algorithm ensured equal treatment by removing human bias, those intentions do not change the reality that the algorithm broke human rights. Because the company had the technical capability to include exceptions for strikes or illness but chose not to, they must bear the responsibility for the discriminatory outcomes.

The Ethics of Responsibility dictates that since the executives were the ones who decided to use the algorithm and reclassify its risk, they are solely accountable for the blindness of the system and its impact on the riders' lives.

Ethics of Institutions

Instead of focusing on individual failings, the Ethics of Institutions examines how the design of the legal and corporate environment influences behavior.

This model highlights that the documentation gap in the EU AI Act—such as **Article 6(3)**—creates an institutional environment where strategic reclassification is easy to perform but difficult to detect. For example, the systemic risk classification for General Purpose AI (GPAI) relies on compute thresholds, yet providers can use clever words to argue that their model does not pose a risk, potentially confusing enforcers who lack deep AI expertise.

To solve this, the institution itself must be redesigned. An ethical institution would move away from black box executive decisions and toward a multi-stakeholder governance model. If Deliveroo had included representatives from trade unions or the riders themselves in the design phase of the Frank algorithm, the discriminatory variables would likely have been identified and corrected before deployment, preventing the dual failure of both law and ethics.

The Proposed Course of Action

After weighing the above-mentioned models, the most ethical course of action is to reject strategic reclassification in favor of honest categorization and proactive, voluntary compliance. While the Economic system might view **Article 6(3)** as a way to escape costs, a long-term ethical perspective suggests that the High-Risk designation should be embraced rather than avoided when human livelihoods are at stake. The company should not wait for a legal mandate or a court verdict to ensure that its algorithm is fair, transparent, and secure. Instead of seeking the minimum legal requirement, the firm should voluntarily

adopt the high-risk standards—such as human oversight and rigorous data governance—even for systems that might technically fall into a grey zone . This proactive stance transforms compliance from a burden into a competitive advantage, as it builds deep trust with employees and customers while insulating the company from the massive financial and reputational risks of future litigation.

Ultimately, the goal of an ethical organization is to ensure that its AI tools are accurate and secure because it is the right way to operate, not simply because the rules say so . By choosing honesty over strategic loophole-seeking, the company addresses the core issue of the Black Box and ensures that its technology serves all stakeholders. This approach requires the executive leadership to bridge the gap between profit and responsibility, moving away from Compliance Theater and toward a model of Explainable AI that respects fundamental human rights. In the long run, this voluntary commitment to high standards creates a stable, sustainable business environment where innovation and ethics are structurally coupled, ensuring that the company thrives without compromising the dignity of those who work under its algorithms.

Application

The Weighted Representation Model

To implement the chosen course of action—proactive and honest classification—the company must move away from isolated executive decision-making. Since including every employee in every discussion is logistically impossible, the company should adopt a Weighted Stakeholder Representation model. In this framework, for every significant AI-related decision, the Impact Radius is calculated.

The weight of a stakeholder's opinion is then assigned based on how directly the decision affects them. For example, in the Deliveroo case, the riders' representatives would hold the highest weight regarding the Reliability Index variables, as their livelihood is directly dependent on that output. This ensures that those with the most at stake have the strongest voice, preventing the intrinsic blindness that occurs when executives prioritize profit codes over human reality.

The Ethics of Disclosure and Consensus

While a full Ethics of Disclosure—where every decision is made collectively by everyone—is time-consuming, a streamlined version can be applied. By pre-categorizing routine business choices, the company can save time, only triggering the full representative discussion when a new algorithm or a significant update to an existing Black Box is proposed.

To facilitate this:

- Categories of stakeholders (e.g., riders, developers, legal experts) vote for representatives to sit on an AI Ethics Committee.
- The remaining employees can communicate concerns to their representatives through internal portals, ensuring their voices are heard without slowing down operations.
- Before a system is registered as Low-Risk, the representative committee must review the procedural task documentation to ensure the company isn't performing Strategic Reclassification.

Continuous Monitoring and Feedback Loops

Application is not a one-time event; it requires constant monitoring. The company must establish Feedback Loops where the algorithm's outputs are regularly audited against the representative committee's ethical standards. If the data shows that the algorithm is again beginning to penalize protected activities (like illness or legal strikes), the representative committee has the power to freeze the system for reassessment. This proactive monitoring ensures the company remains in the Profit/Legal quadrant of structural coupling, avoiding the Loss/Illegal outcomes of court-mandated corrections.

Conclusion

The Deliveroo Frank algorithm case serves as a definitive case study for the necessity of the EU AI Act's risk-based framework. It illustrates that when a company prioritizes infinite scalability and economic efficiency through black box logic, it inherently risks creating a system that is blind to fundamental human rights. The Court of Bologna's ruling proved that technical neutrality is not a defense for discriminatory outcomes, especially when a provider has the technical means to include ethical exceptions but chooses not to for the sake of profit.

As the EU AI Act moves into full enforcement, the temptation for strategic reclassification under **Article 6(3)** remains a significant GRC challenge. However, as analyzed through the lenses of Utilitarianism, Deontology, and Social Systems Theory, bypassing high-risk obligations is a short-term strategy that leads to long-term legal and reputational failure. The documentation gap in the Act must be bridged by a corporate commitment to Ethics of Responsibility, where leaders are held accountable for the impact of their tools regardless of their technical intent.

Ultimately, the transition from a blind algorithm to a trustworthy one requires a move toward Inclusive Governance. By adopting a weighted representative model, companies can ensure that the voices of those most impacted—such as the riders—are integrated into the decision-making process. Moving forward, the goal for any AI-driven enterprise should be proactive and voluntary compliance. By embracing the rigorous standards of the High-Risk category even in grey zones, a company moves beyond compliance theater and toward a sustainable model where technological innovation and human dignity are structurally coupled.

References

- **Official Legislation:** European Parliament & Council of the European Union. (2024). *The EU Artificial Intelligence Act*. <https://artificialintelligenceact.eu/>
- **Academic Journal (The Springer Article):** Meijerink, J., Boons, M., Keegan, A., & Marler, J. (2022). Algorithmic Management: Principles, Algorithms, and Programs. *Business & Information Systems Engineering*, 64(1), 1–10. <https://link.springer.com/article/10.1007/s12599-022-00764-w>
- **Legal Analysis (Clifford Chance):** Clifford Chance. (2021). *The Italian Courts Lead the Way on Explainable AI*. Talking Tech. <https://www.cliffordchance.com/insights/resources/blogs/talking-tech/en/articles/2021/06/the-italian-courts-lead-the-way-on-explainable-ai.html>
- **News/Human Rights Report:** Business & Human Rights Resource Centre. (2021). *Italy: Court rules Deliveroo used discriminatory algorithm for ranking riders*. <https://www.business-humanrights.org/en/latest-news/court-rules-deliveroo-used-discriminatory-algorithm/>
- **Course Material:** Prof. Dr. Bernd Banke. (2025). *Business Ethics - A contradiction in itself?* [Class notes]. Hochschule Karlsruhe.